



Vector, SIMD, & GPUs

Module 10

Module Syllabus

- Vector processors
- Single instruction, multiple data (SIMD) architectures
- SIMD case study
- Graphics processing units (GPUs)
- GPU case study

Motivation

- Multicore processors and multithreading take advantage of thread-level parallelism.
 - Allows us to exploit thread-level parallelism (within a program) and process-level parallelism (across applications)
- They are relatively simple extensions to a conventional core (at least conceptually).
 - For multicore, duplicate the core several times on the chip and add extra logic (e.g., for coherence).
 - For multithreading, add storage for multiple thread contexts and associated pipeline changes.
- However, other architectural choices allow us to extract different forms of parallelism.
 - Such as data-level parallelism
- This module studies these architectures compared to a generic multicore.

Vector Processors


Vector Processors

- Vector processors explicitly exploit data-level parallelism.
 - When instructions are applied to multiple independent data items
- Vector processors gather data from memory into large (vector) registers.
 - Then, conceptually perform operations on these items together
 - Essentially performing many register-register operations with the same opcode
 - Results are scattered back out to memory.
- For vectorizable data, vector processors:
 - Provide energy-efficient computation (amortizing the costs of fetch and decode).
 - Hide memory latencies through pipelining operations.

Example Executing Code on a Vector Processor


Contiguous data in memory

```
for (int i=0; i<64; ++i) {  
    A[i] = B[i] + C[i]  
}
```



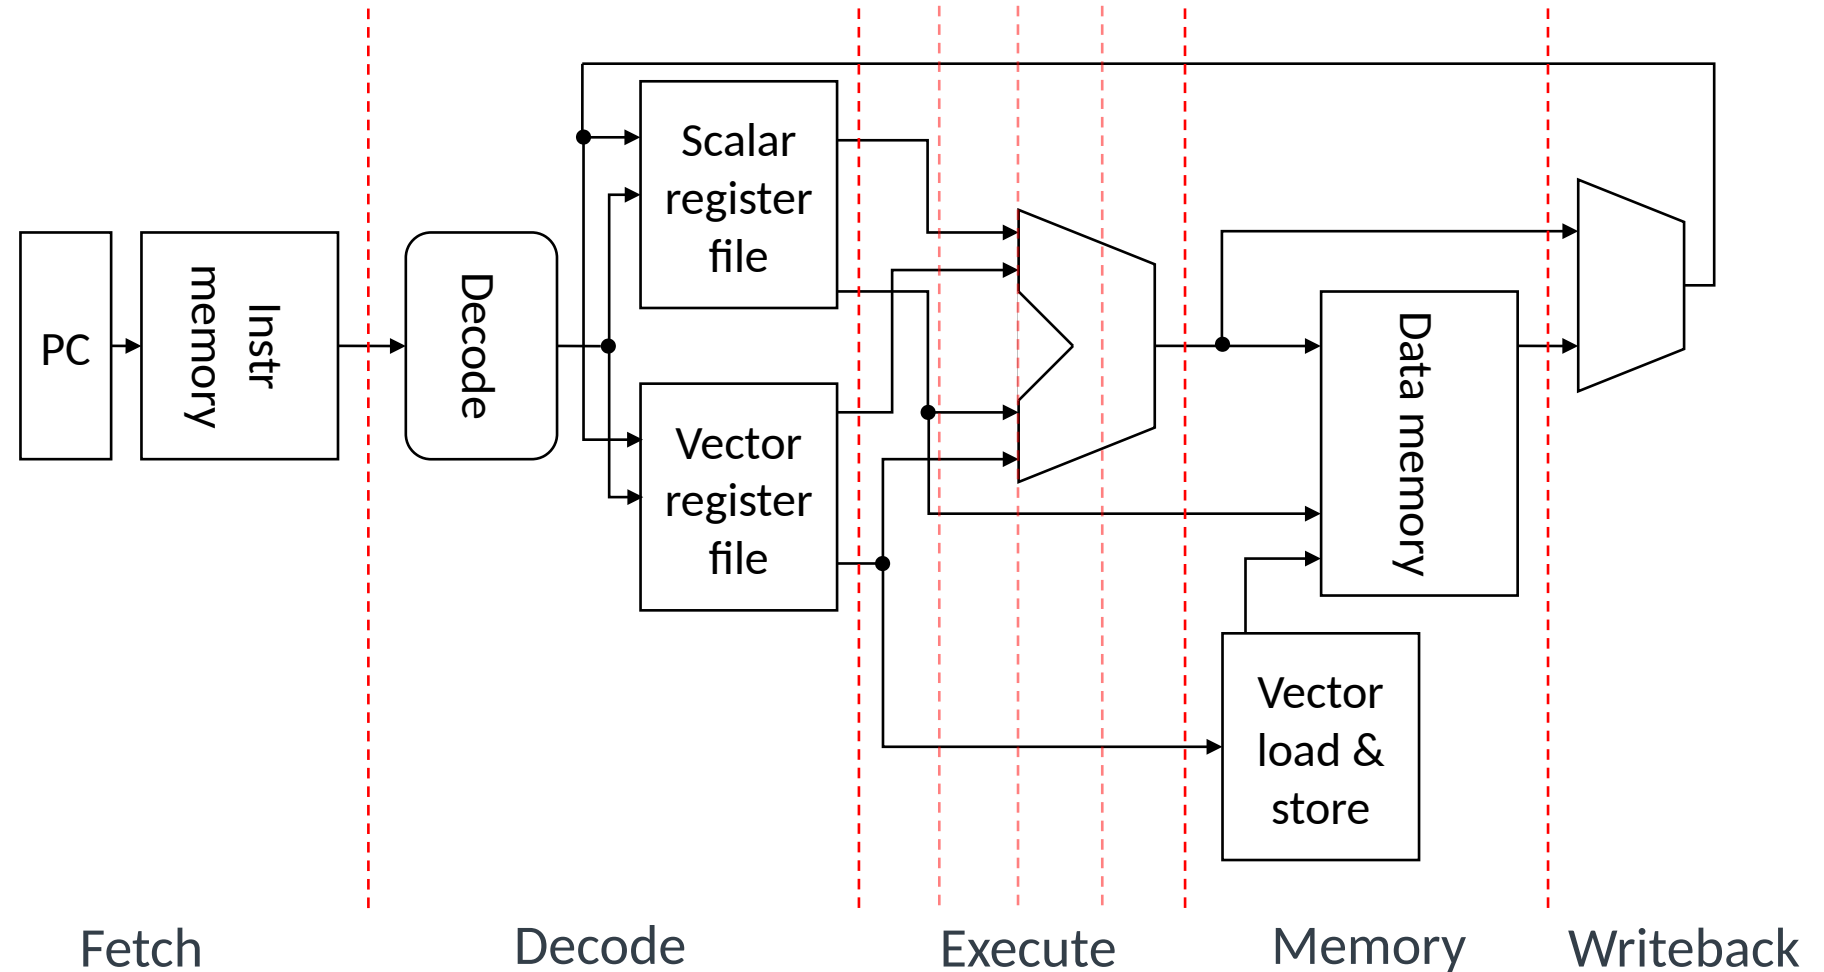
Dispersed data in memory

```
for (int i=0; i<64; ++i) {  
    A[i] = B[D[i]] + C[i]  
}
```



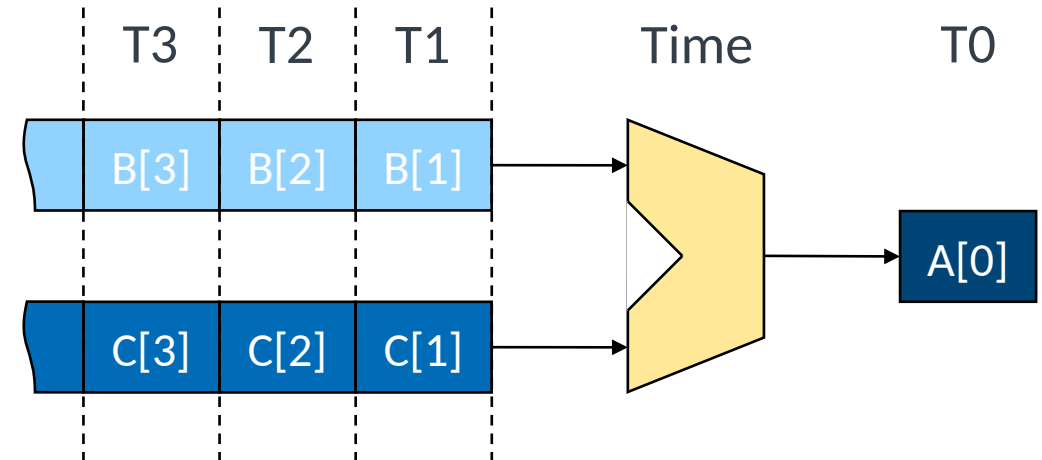
An Example Vector Processor

- Vector functional units fed by
 - Vector registers provide independent data to operate on.
 - Scalar registers provide additional data or memory addresses.
- Vector FUs are fully pipelined.
 - To start processing a new element on each clock cycle



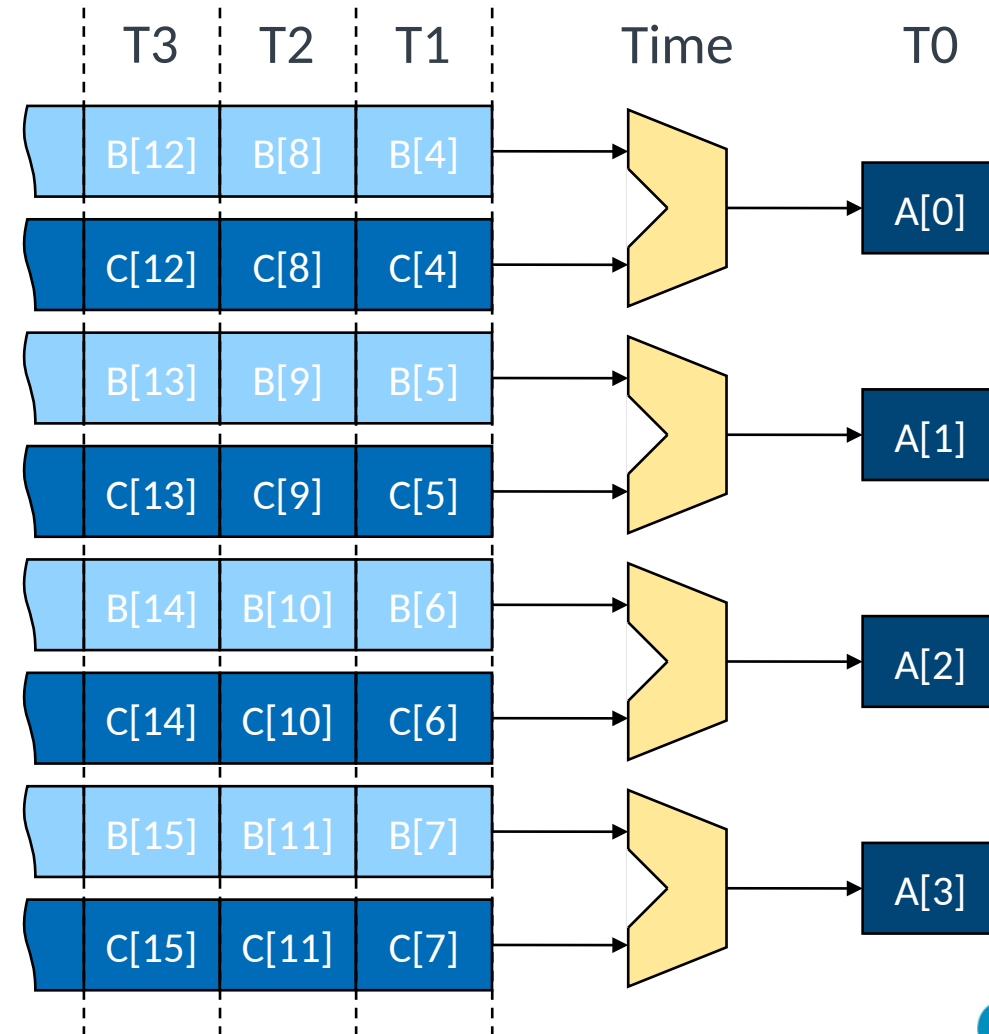
Instruction Execution

- The core operates on vectors of data
 - However, this doesn't mean that the whole vector must be processed at once.
- This is a simple vector functional unit.
 - It can only process one element at a time.
- Operation on a new data element starts on each clock cycle.
 - The FU is pipelined to support multi-cycle operations.
- The ALU needs as many cycles as there are vector elements to process all data.
 - Plus the time to perform one operation



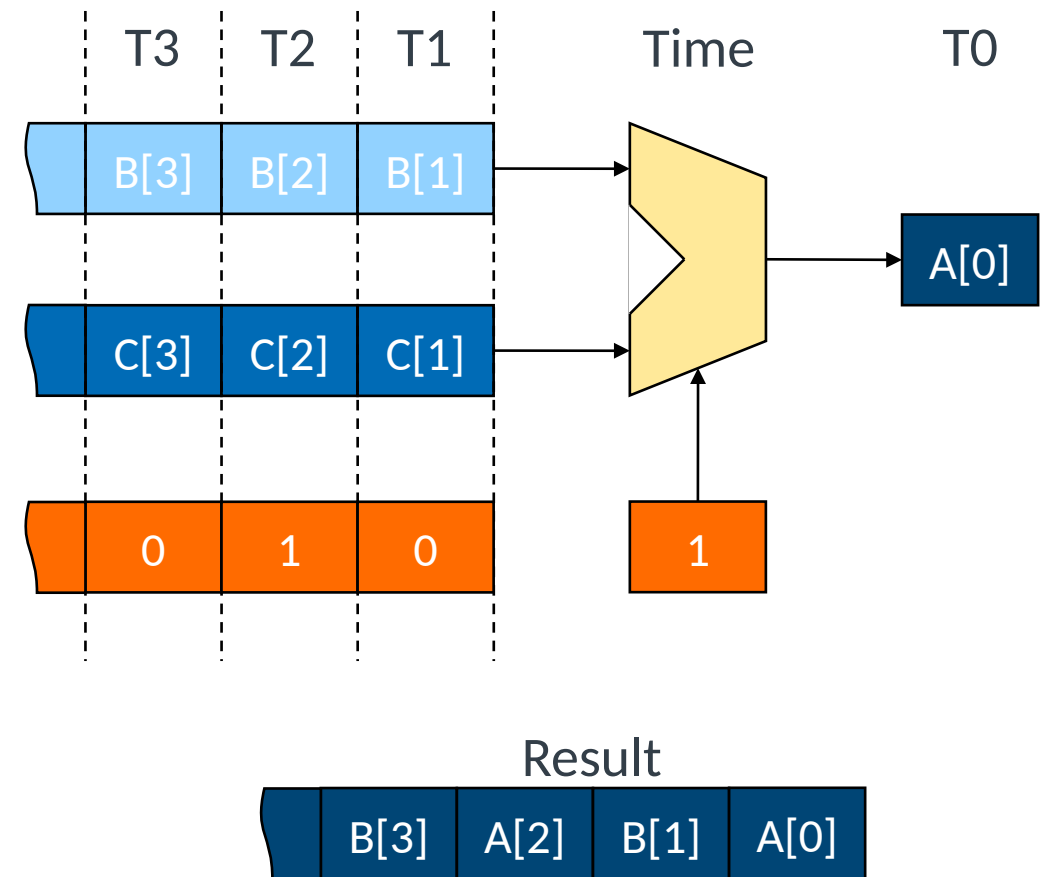
Instruction Execution

- Duplicating the functional units increases performance.
 - Since we now start execution on multiple data elements on each clock cycle
- To help this, we can partition the vector register file into lanes.
 - With one functional unit of each type per lane
 - Data elements are interleaved across lanes to produce data in the correct order.



Predication

- Sometimes you don't want to do computation across the whole vector.
 - Only certain elements of the computation should be performed.
- Predication is one method for achieving this.
- Include one or more predicate registers.
 - Usually, these contain one bit per vector element.
- Bits from the predicate register say whether the operation on the corresponding elements goes ahead.
 - If not, some other value is placed in the destination.
 - E.g., forward the value from the first source



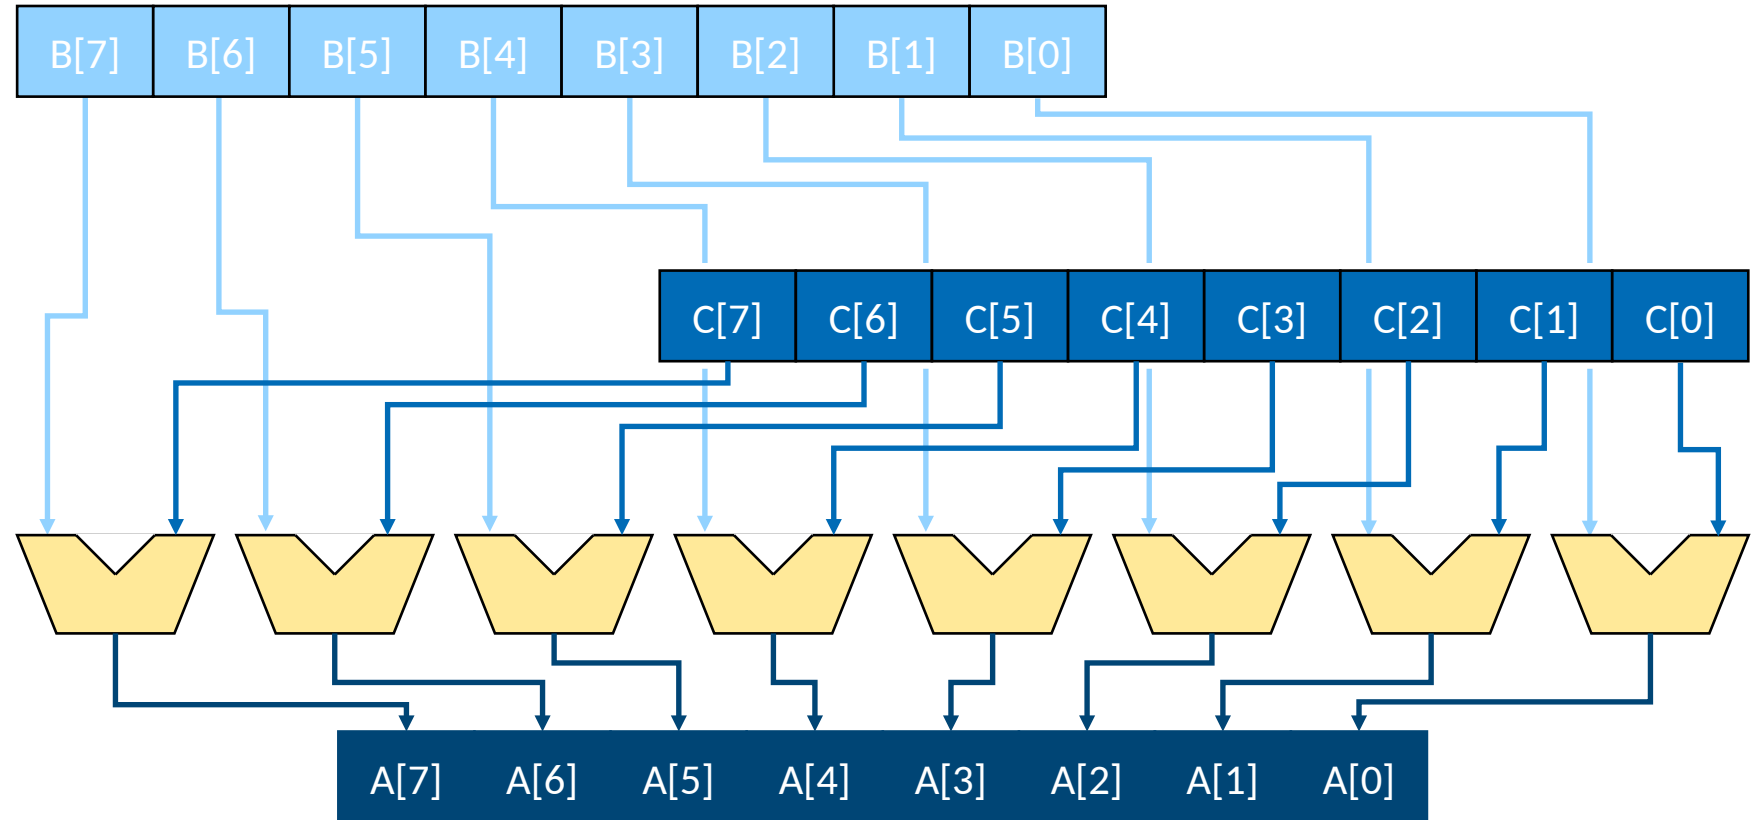
Single Instruction, Multiple Data (SIMD)

SIMD

- SIMD execution brings a (constrained) form of vector processing to general-purpose CPUs.
 - The key idea remains to exploit data-level parallelism.
 - And obtain performance benefits (with energy-efficiency benefits, too, if possible)
- SIMD also provides more efficient processing for certain codes.
 - E.g., multimedia workloads, such as image recognition and object detection, where operations are on pixel color values that are 8 or 16 bits in length
 - Scalar execution would put each value in its own 32-bit register.
 - SIMD packs them into a vector register (e.g., 256-bit vector of $16 * 16$ -bit data elements).
 - Then operates on each element independently
- However, SIMD typically operates on all elements concurrently.
 - Contrast with vector processing in the previous slides.

SIMD Execution

- In SIMD, operations on all data elements occur at the same time.
- The processor provides as many FUs as there are data elements in a vector.



SIMD vs Vector Processing

Vector processing


- Operations produce data for elements of the output vector over multiple cycles.
- Maximum vector length can change with each processor generation without requiring ISA changes.
- Best performance when the majority of the code is vectorizable – the scalar architecture is relatively simple.

SIMD

- Operations produce data for elements of the output vector in the same cycle.
- In most SIMD ISAs, the vector length is hardcoded, so increasing it requires new instructions.
- Vector operations are used to increase performance and generally augment a high-performance scalar architecture.

Case Study: AArch64 NEON

- NEON is a wide SIMD architecture developed for multimedia applications.
- Registers are considered as vectors of elements of the same data type (128 bits in size).
 - Integer signed and unsigned 8-bit, 16-bit, 32-bit, 64-bit
 - Floating point half, single and double precision
 - Instructions usually perform the same operation in all lanes.





Case Study: AArch64 SIMD Instruction Types

Type	Examples
Arithmetic	ADD, SUB, MUL, MLS, SMIN, SMAX
Saturating math	UQADD, UQRSHL, SQDMULL
Narrowing instructions	SUBHN, ADDHN, RSUBHN
Widening instructions	SSUBL, UMULL2, UABDL2
Integer compare	CMGT, CMHS, CMTST
Logical operations	ORR, AND, BIC, EOR
Floating point operations	FADD, FABD, FDIV, FRINTZ
Data movement	DUP, INS, MOV, SMOV, UMOV
Load/store instructions	LDR, LDP, LD1, ST1, LD4R

More information at <https://developer.arm.com/architectures/cpu-architecture/a-profile/exploration-tools>

Case Study: AArch64 SIMD Register Bank

- Separate set of 32 registers, 128 bits wide, **V0** - **V31**
- For access to a scalar
 - **Qn** to access 128-bit data (in **Vn[127:0]**)
 - **Dn** to access 64-bit data (in **Vn[63:0]**)
 - **Sn** to access 32-bit data (in **Vn[31:0]**)
 - **Hn** to access 16-bit data (in **Vn[15:0]**)
 - **Bn** to access 8-bit data (in **Vn[7:0]**)



Case Study: AArch64 SIMD Data Types and Sizes


- Data types
 - Unsigned integer (U8, U16, U32, U64) and signed integer (S8, S16, S32, S64)
 - Integer of unspecified sign (I8, I16, I32, I64)
 - Floating point number (F16, F32, F64) and polynomial (P8, P16)
- Data sizes
 - A single scalar value of a floating-point or integer type (Q, D, S, H, B)
 - A 64-bit wide vector containing two or more elements (2S, 4H, 8B)
 - A 128-bit wide vector containing two or more elements (2D, 4S, 8H, 16B)
- Data type is specified as the instruction prefix; size is specified in the operands' type sizes.
 - For example, FADD V2.4S, V0.4S, V1.4S
 - Perform a floating point add of four single-precision values in registers V0 and V1, putting the result in V2
 - But not all combinations of data type size and operation are available (see the Arm Architecture Reference Manual for valid combinations).

Case Study: Scalable Vector Extension (SVE)

- SVE for Armv8-A Arch64
 - Next-generation SIMD instruction set for AArch64
- Motivated by a need for better vectorization of “real-world” applications
 - Enabling different CPUs to implement different vector lengths while sharing a common ISA
 - As a result, a program written for one CPU should work on another with no changes.
 - Add support for data-set lengths that are not a multiple of the vector width
 - Also aims to reduce the initial porting effort to use a new ISA, scalable for future designs
- A vector-length-agnostic architecture
 - Implementations can range from 128 bits up to 2048 bits.
 - Introduces new vectorization techniques

Case Study: SVE Vectors and Predicates

- 32 new scalable vector registers (**Z0-Z31**), length determined by the implementation
 - Bottom 128 bits overlay the floating-point & NEON vector register bank (**V0-V31**); top bits zeroed on a write



- 16 new scalable predicate registers (**P0-P15**) ; this example is for 256-bit vector registers.
 - Have 1/8th of a vector register's length: 1 bit of predicate register is mapped to 1 byte of vector register

Zx	8-bit	8-bit	8-bit	8-bit	...	8-bit	8-bit	8-bit	8-bit	8-bit element
Px	1	1	1	1	...	1	1	1	1	32 bits
Zx	16-bit		-		...	16-bit		-		Unpacked 16-bit element
Px		1		0	...		1		0	32 bits

GPUs

GPU Evolution

- GPUs originally developed as specialized hardware for graphics computation
 - Now used as programmable accelerators for highly data-parallel workloads
- GPU hardware evolution closely tied to evolution in usage patterns
 - Desire for improved visual effects driven by games
 - More realism, more effects, more screen resolution, more frames per second
- Originally a fixed-function pipeline, programmability was added gradually to many stages.
 - Now the bulk of the GPU is a programmable data-parallel architecture.
 - Some fixed-function hardware remains for graphics.
 - New hardware being added to cope with modern workloads, e.g., machine learning

GPU Design Principles

- CPU design is about making a single thread run as fast as possible.
 - Pipeline stalls and memory accesses are expensive in terms of latency.
 - So increased logic was added to reduce the probability/cost of stalls.
 - Use of large cache memories to avoid memory misses
- GPU design is about maximizing computation throughput.
 - Individual thread latency not considered important
 - GPUs avoid much of the complex CPU pipeline logic for extracting ILP.
 - Instead, each thread executes on a relatively simple core with performance obtained through parallelism.
 - Single instruction, multiple threads
- Computation hides memory and pipeline latencies.
- Wide and fast bandwidth-optimized memory systems

SIMT vs SIMD

SIMT

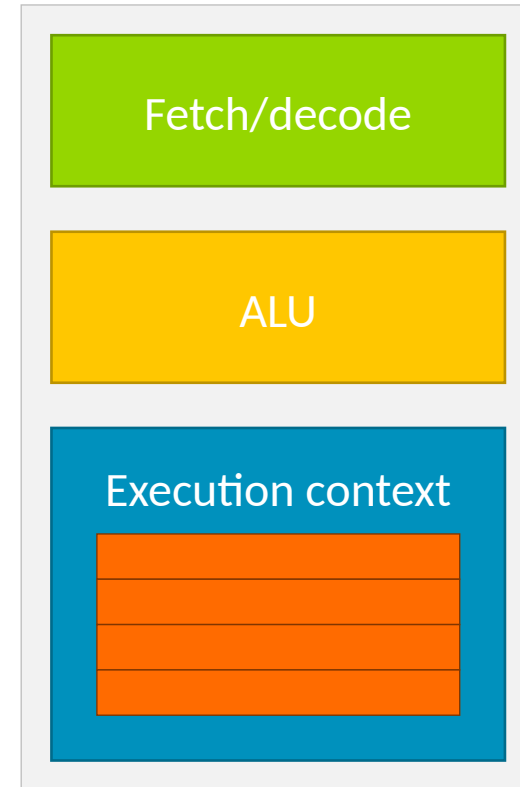
- Single instruction, multiple threads
- Takes advantage of data-level parallelism
- Can be considered a constrained form of multithreading
- Many threads each with their own state
 - Operating on scalar registers
 - With their own local memory

SIMD

- Single instruction, multiple data
- Takes advantage of data-level parallelism
- Can be considered a constrained form of vector processing
- One thread operating on vector registers

One Processing Element

- A single thread runs on a simple processing element.
- Short pipeline
- The execution context consists of the thread's state.
 - E.g., registers and local memory
- But fetch and decode are costly.



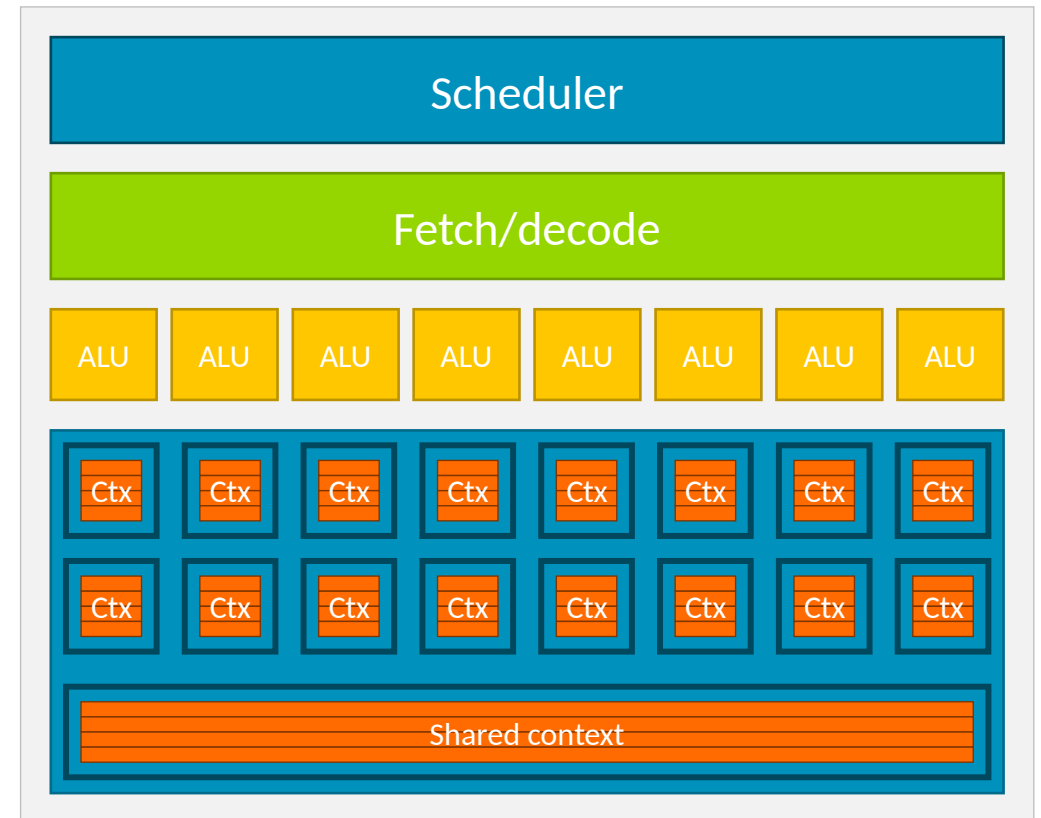
Multiple Processing Elements

- SIMT execution of threads
- Cost of fetch and decode spread across all threads in a work group (OpenCL terminology)
- Each thread has its own context.
 - And some shared context
- Multiple functional units for parallelism
 - One per thread for cheap units (e.g., simple ALU)
 - Fewer expensive units than threads (e.g., sqrt)
- However, stalls are costly.




Minimizing Stalls

- Include support for multiple work groups
 - Each is independent of all others.
 - And this is guaranteed by the compiler.
 - Which means there is no fixed ordering of groups.
- A scheduler chooses work groups to run.
 - Maintains a list of ready work groups
 - Makes a choice each cycle
 - Some GPUs can schedule more than one work group per cycle.
 - Hides latency when a work group stalls
- This system is sometimes called a shader core.



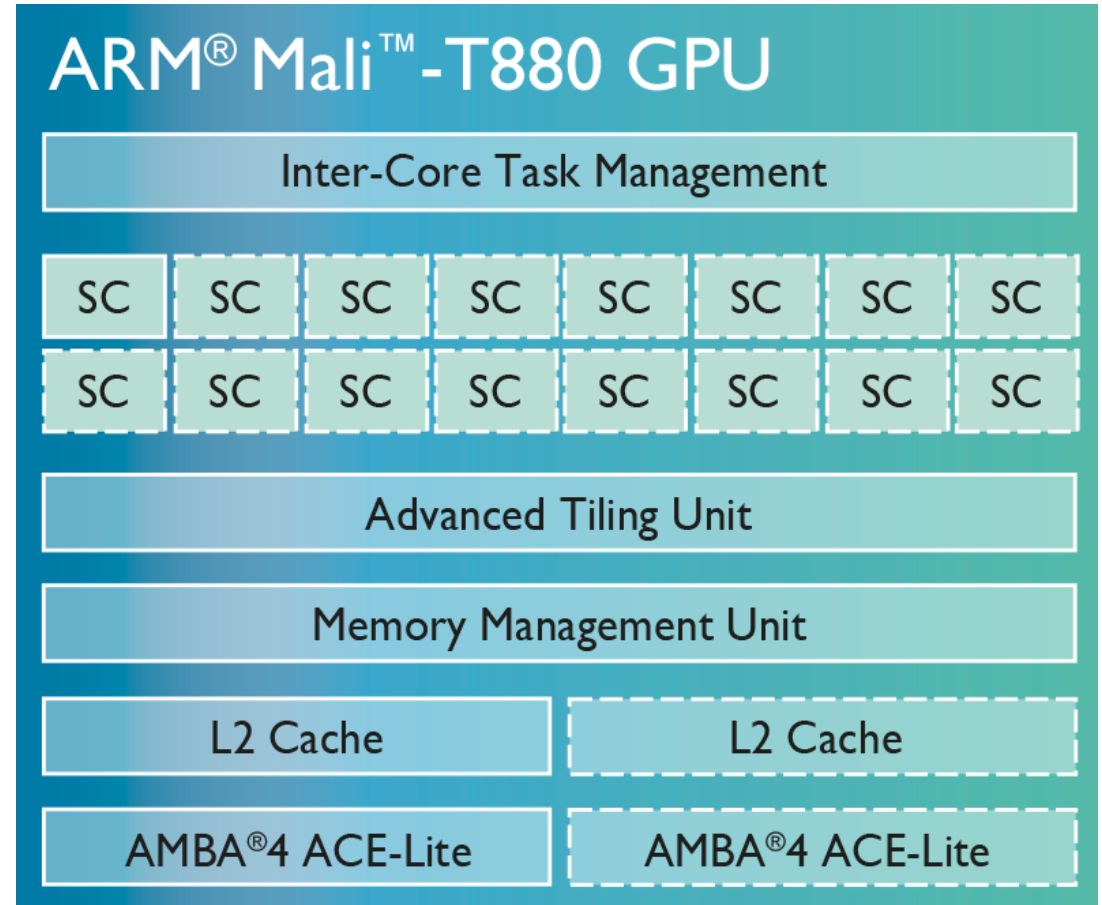
Scaling Out

- Multiple instances of each shader core provided together
 - Each independent of the others
 - Each processes a subset of the work groups
- Massively increases parallelism
- Memory hierarchy provided, too
 - Shared L2 cache reduces memory bandwidth requirements.
 - Smaller caches local to each multithreaded core



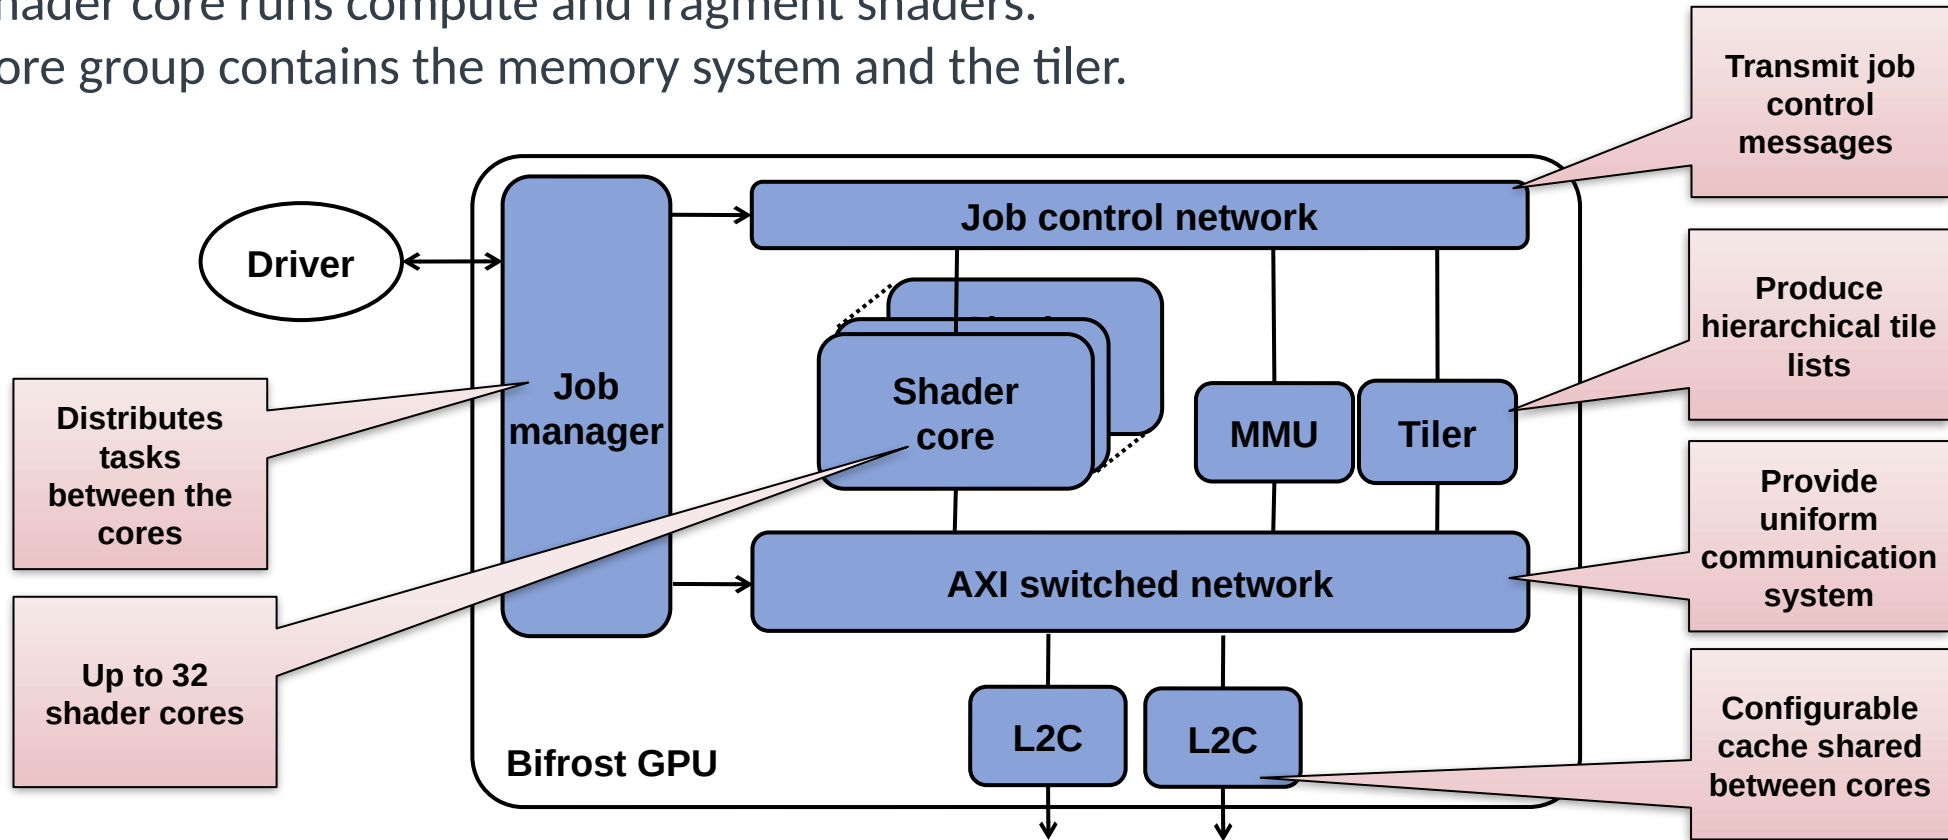
Case Study: Mali T880 GPU

- Up to 16 shader cores (SC)
 - Each core supports multiple threads and operations.
- Performance
 - 30.6G FLOPS at 900 MHz
- API support
 - OpenGL ES 1.1, 1.2, 2.0, 3.0, 3.1
 - OpenCL 1.1, 1.2
 - DirectX 11 FL11_2
 - RenderScript
- Usage in SoCs
 - Exynos 8890, Helio X20 (MT6797), Kirin 950
- Used for both graphics processing and high-performance computing




Case Study: Arm Mali Bifrost GPU Architecture

- Mali Bifrost GPU consists of 3 main blocks or groups.
 - The job manager interacts with the driver and controls the GPU HW.
 - The shader core runs compute and fragment shaders.
 - The core group contains the memory system and the tiler.



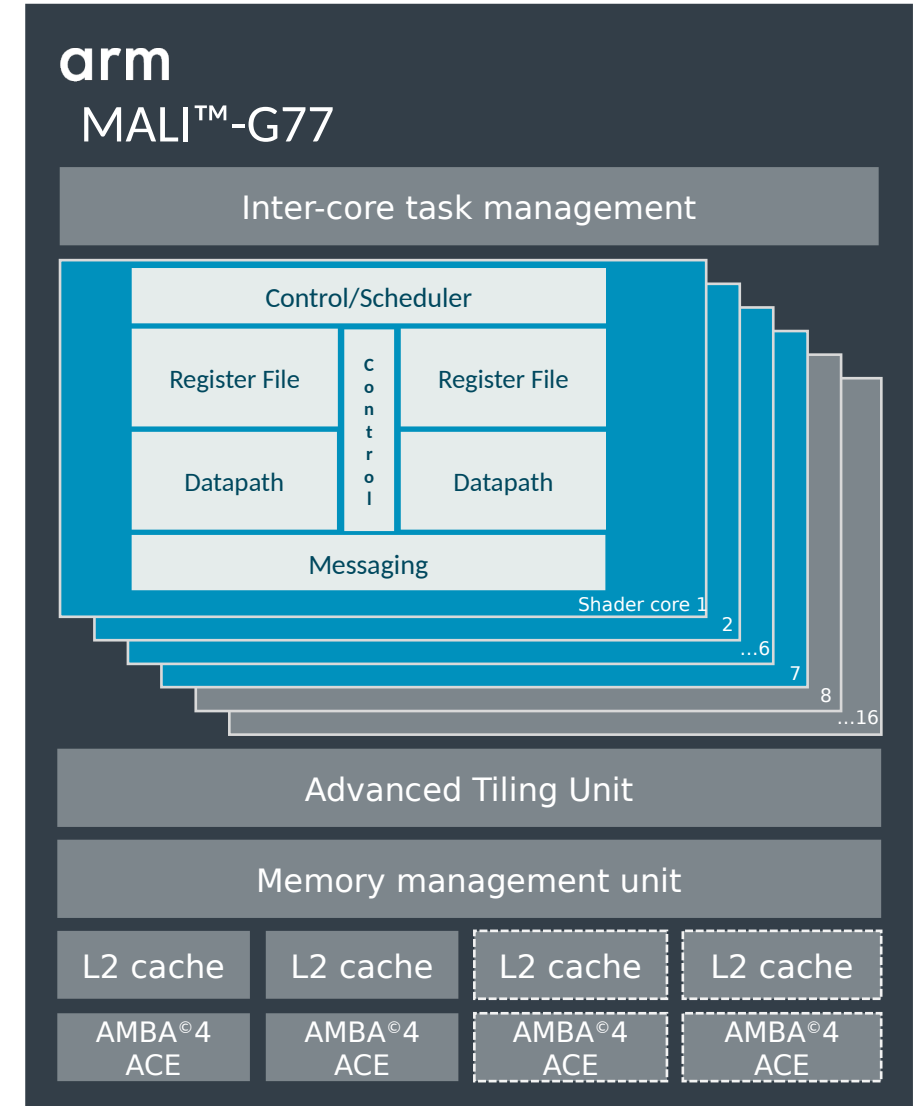
Case Study: Arm Mali-G76 GPU

- Third generation of the Bifrost architecture
- Maximum 20 shader cores (SC)
 - Wider execution engines with double the number of lanes
- Performance
 - Complex graphics and machine-learning workloads
- API support
 - OpenGL ES 1.1, 2.0, 3.1, 3.2
 - OpenCL 1.1, 1.2, 2.0 Full profile
 - Vulkan 1.1
- Shared L2 cache with 2 or 4 slices



Case Study: Arm Mali-G77 GPU

- First generation of the Valhall architecture
 - Warp-based execution model
 - New instruction set with operational-equivalence to Bifrost
 - Dynamic instruction scheduling in hardware
- Configurable 7 to 16 shader cores
- Single execution engine per shader core
- Configurable 2 to 4 slices of L2 cache
 - 512 KB to 4 MB in total
- Texture mapper – 4 texture element (texel) per cycle
- Supports Vulkan and OpenCL



Summary

- General-purpose CPUs optimize scalar single-threaded performance
 - But there are many domains where data-level parallelism is common.
 - Other architectures can exploit this efficiently.
- Vector processing is highly efficient for regular workloads computing on arrays of values.
 - The basis of early supercomputers
- SIMD processing brings a form of this to the general-purpose CPU.
 - Useful for multimedia codes and data-parallel operations
- GPUs exploit a different form, SIMT, with massive parallelism.
 - Simple cores but multithreading and multiprocessing hide latency